



A shared clinical graph, not a flat patient record

One typed graph per decision-time visit: patient, visits, complaints, labs and vitals, with numeric change between earlier and index results kept on the edge. Saved as a reusable project asset.



CEI-GNN: a diagnosis score you can decompose

Each class score is an exact sum of signed evidence terms: positive supports the class, negative opposes it.



PLE = piecewise-linear encoding of numeric values, fitted on training data (K = 4 selected). CEI v3 builds on the additive v2 pair mode; the multiplicative pair interaction was not successful.

What the evidence supports

Primary screen, macro-F1 (3-seed mean)

Method	Macro-F1
CEI-GNN v3 (arm C)	0.6418
CEI-GNN v2 (arm A)	0.6420
ProtGNN	0.6371
XGBoost	0.6371

No decisive screen gain yet

No contrast passes the preregistered rule; v3 over v2: -0.0002, 95% CI [-0.0047, +0.0042].

What we can say

Interpretable by construction. Every CEI score splits exactly into signed evidence terms.

Explanations are model-faithful, not clinical proof. Against Grad and IG on 500 screen visits, CEI is favoured on fidelity-minus and sparsity, not uniformly on fidelity-plus; ProtGNN has no native explanation to compare.

Accuracy is on par, not superior. Differences are within about 0.005; no winner claim.