



A shared clinical graph, not a flat patient record

One typed graph per decision-time visit: patient, visits, complaints, labs and vitals, with numeric change between earlier and index results kept on the edge. Saved as a reusable project asset.



CEI-GNN: a diagnosis score you can decompose

Each class score is an exact sum of signed evidence terms: positive supports the class, negative opposes it.



PLE = piecewise-linear encoding of numeric values, fitted on training data (K = 4 selected). CEI v3 builds on the additive v2 pair mode; the multiplicative pair interaction was not successful.

What the evidence supports

Primary screen, macro-F1 (3-seed mean)

Method	Macro-F1
CEI-GNN v3 (arm C)	0.6418
CEI-GNN v2 (arm A)	0.6420
ProtGNN	0.6371
XGBoost	0.6371

No decisive screen gain yet

No contrast passes the preregistered rule; v3 over v2: -0.0002 , 95% CI $[-0.0047, +0.0042]$.

What we can say

Interpretable by construction. Every CEI score splits exactly into signed evidence terms.

Explanation evidence is mixed. Using Grad and IG on 500 screen visits, CEI outperforms ProtGNN on fidelity-minus and sparsity, not uniformly on fidelity-plus. Native ProtGNN attribution was unavailable in this implementation.

Superiority is not established. The screen is inconclusive; this does not prove equivalence.